

CS3510

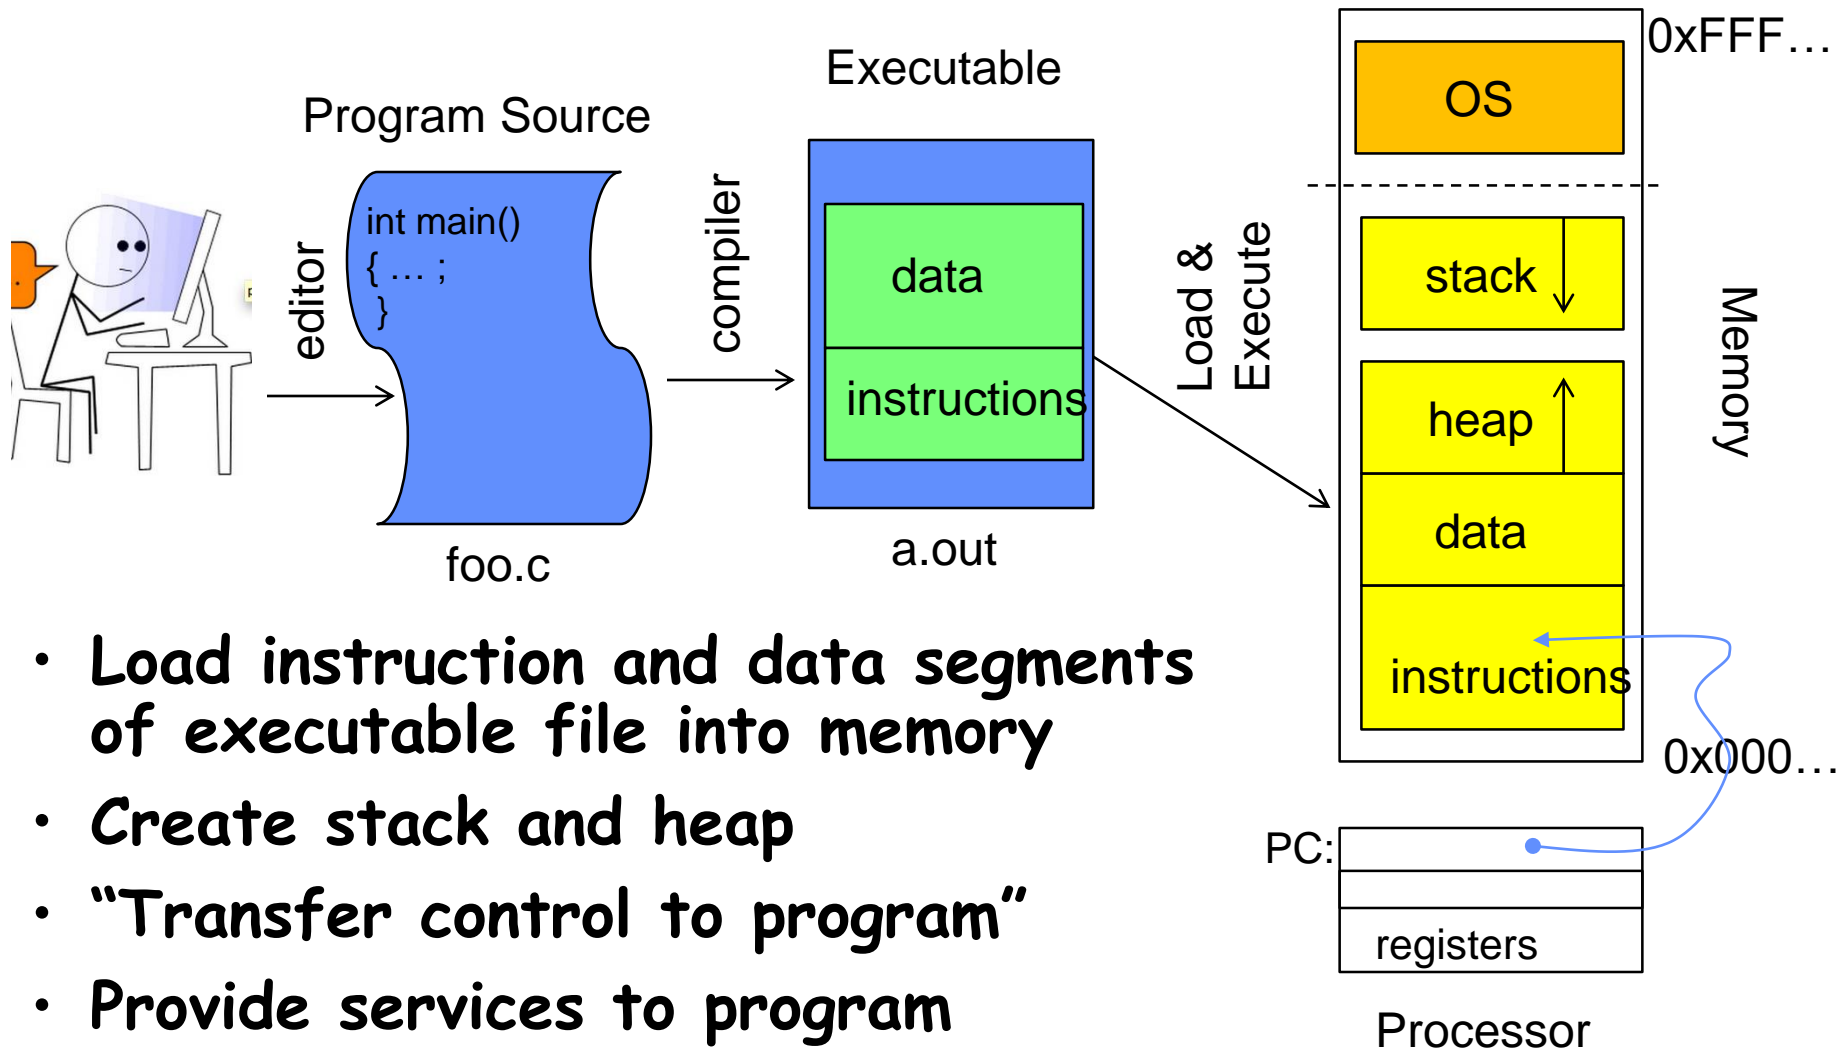
Operating Systems

Processes

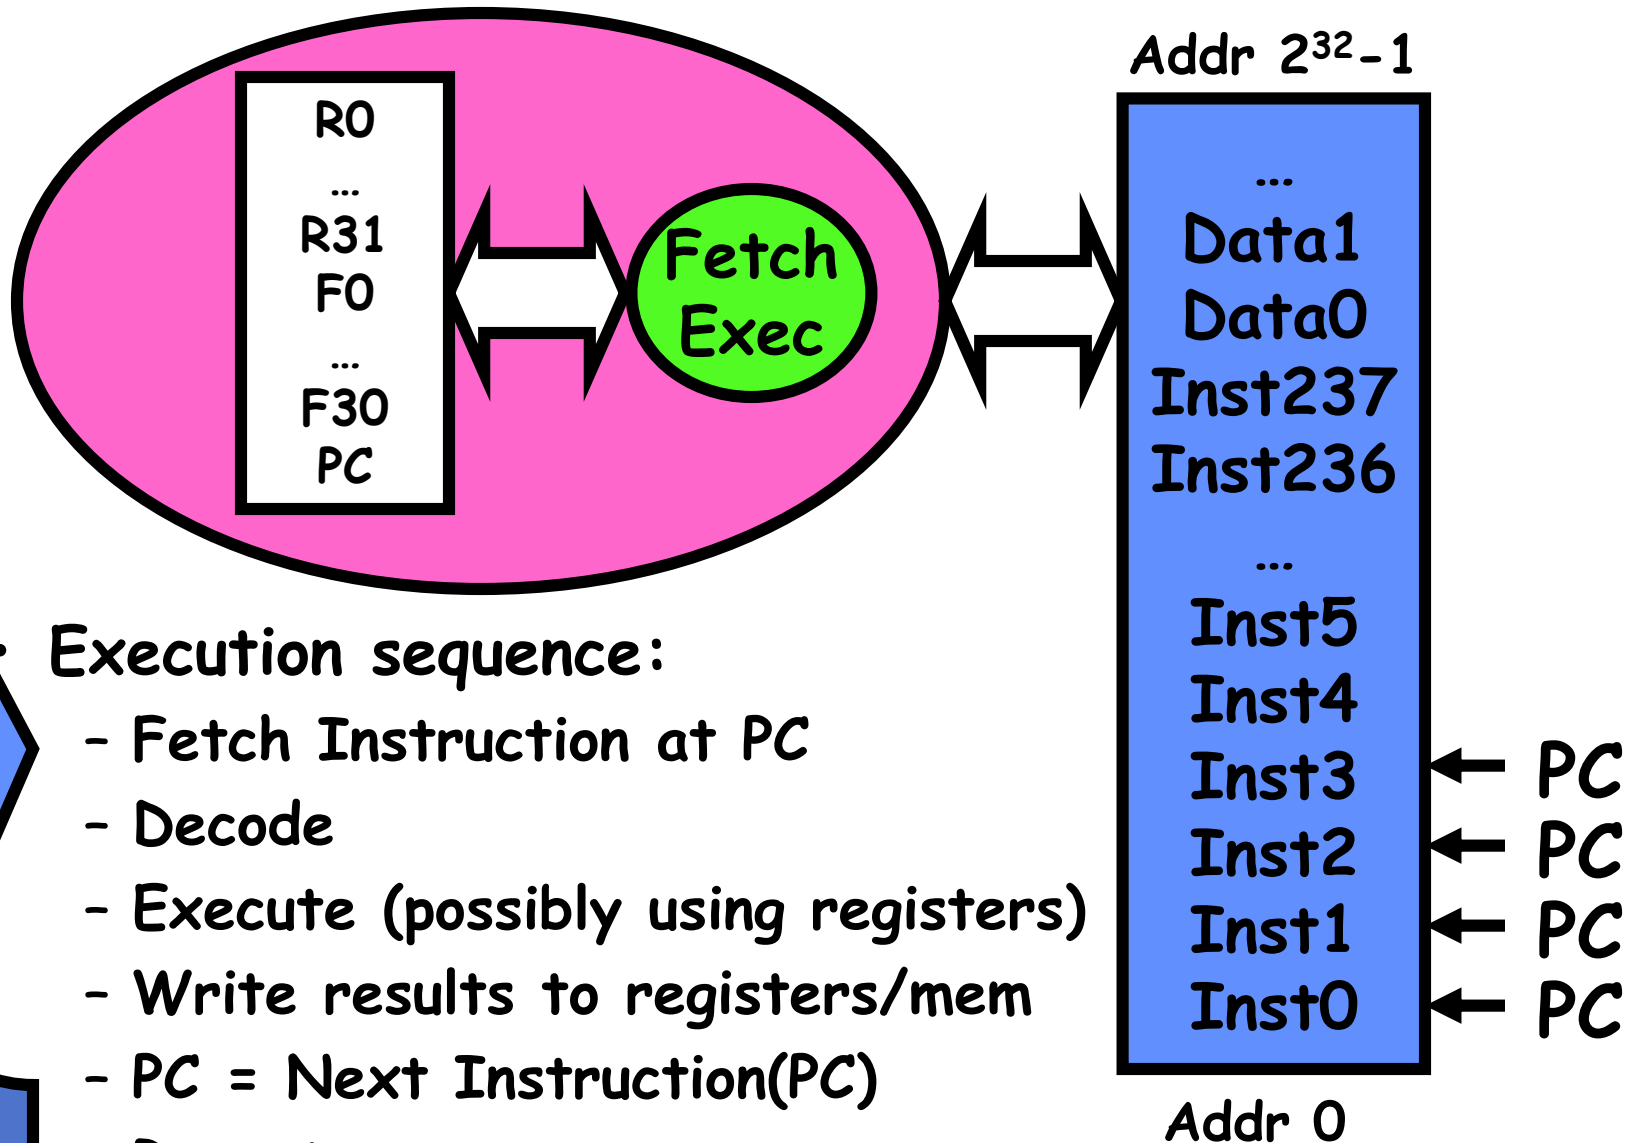
Bheemarjuna Reddy Tamma
IIT HYD

Note: Some slides and/or pictures in the following are adapted from slides ©2005 Silberschatz, Galvin, and Gagne. Many slides are from Prof John Kubi, UCB.

OS Bottom Line: Run Programs



What happens during execution?



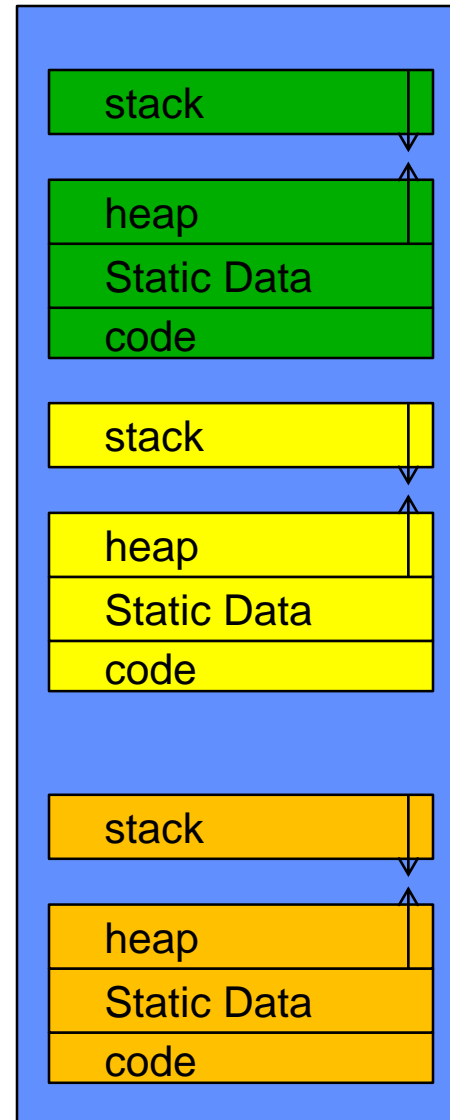

Thread of Control

- **Thread: Single unique execution context**
 - Independent Fetch/Decode/Execute loop
 - Operating in some Address space
 - **Program Counter, Registers, Execution Flags, Stack**
- Certain registers hold the *context* of thread
 - PC register holds the address of executing instruction in the thread
 - Stack pointer holds the address of the top of stack
 - May be defined by the instruction set architecture or by compiler conventions
- A thread is executing on a processor when it is **resident** in the processor registers
 - Registers hold the root state of the thread
 - » The rest is “in memory”

Concurrency vs Parallelism

- Uniprogramming: *one thread at a time in the system*
 - MS/DOS, early Macintosh, Batch processing
 - Easier for operating system builder
 - Does this make sense for personal computers?
- Multiprogramming: *more than one thread at a time in the system*
 - Multics, UNIX/Linux, OS/2, Windows 10, Mac OS X
 - Often called “multitasking”, but multitasking (aka time sharing and concurrency) is bit different from it
 - » CPU executes multiple jobs alternatively by switching from one to other, but switching is so frequent to provide interactive computing (time slice)
- Multi-processor (or ManyCore) System ⇒ Multiprogramming or Multitasking?
 - No, it's parallel programming!


Multiprogramming/Concurrency-Multiple Threads of Control



Traditional UNIX Process

- **Process:** *Operating system abstraction to represent what is needed to run a single program*
 - Process is an instance of a program in execution
 - Often called a "HeavyWeight Process"
 - There is no concurrency in a heavyweight process → a single thread!
 - Collection of data structures to fully describes how far the execution of the program has progressed
 - Formally: a single, sequential stream of execution in its own address space
- **Process has two parts:**
 1. **Sequential Program Execution Stream**
 - » Code executed as a *single, sequential* stream of execution
 - » Includes State of CPU registers
 2. **Protected Resources:**
 - » Main Memory State (contents of Address Space)
 - » I/O state (i.e. file descriptors)

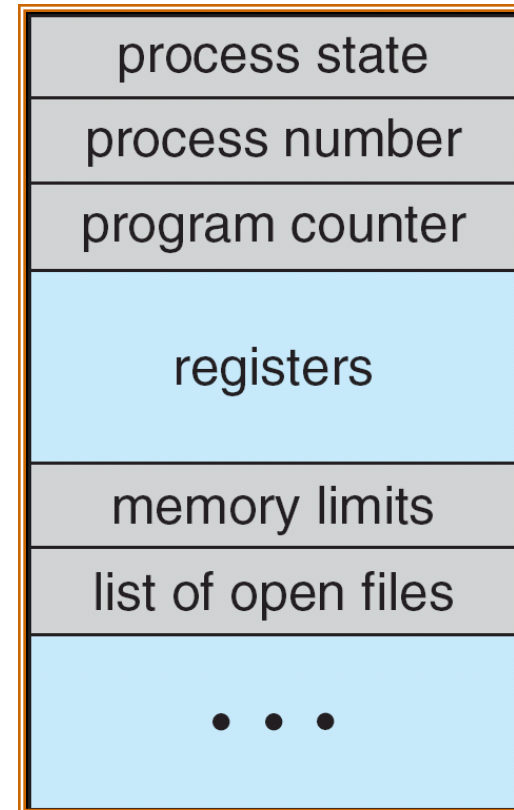
Process = Program ??



- Program is **passive** entity, but process is **active** entity
 - Program becomes process when its executable file is loaded into memory → click on icon or type it on command line (shell)
- More to a process than just a program:
 - Program is just a part of the process state
 - I run **vi editor** on lectures.txt, you run it on homework.c - Same program, but different processes
- Less to a process than a program:
 - A program can invoke more than one process
 - cc starts up cpp, cc1, cc2, as, and ld

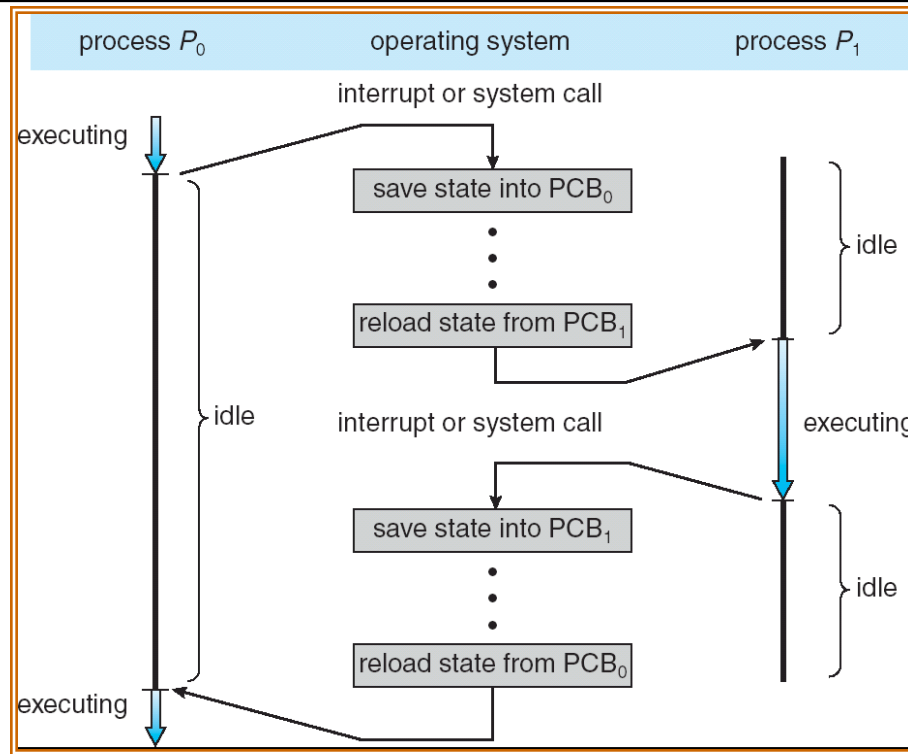
How do we multiplex processes?

- The current state of process held in a DS called process control block (PCB):
 - This is a “snapshot” of the execution and protection environment
 - Only one PCB active at a time
- Give out CPU **time slice** to different processes (**Scheduling**):
 - Only one process “running” at a time
 - Give more time to important processes e.g, I/O, foreground jobs
- Give pieces of resources to different processes (**Protection**):
 - Controlled access to non-CPU resources
 - Sample mechanisms:
 - » Memory Mapping: Give each process its own (virtual) address space
 - » Kernel/User duality: Arbitrary multiplexing of I/O through system calls



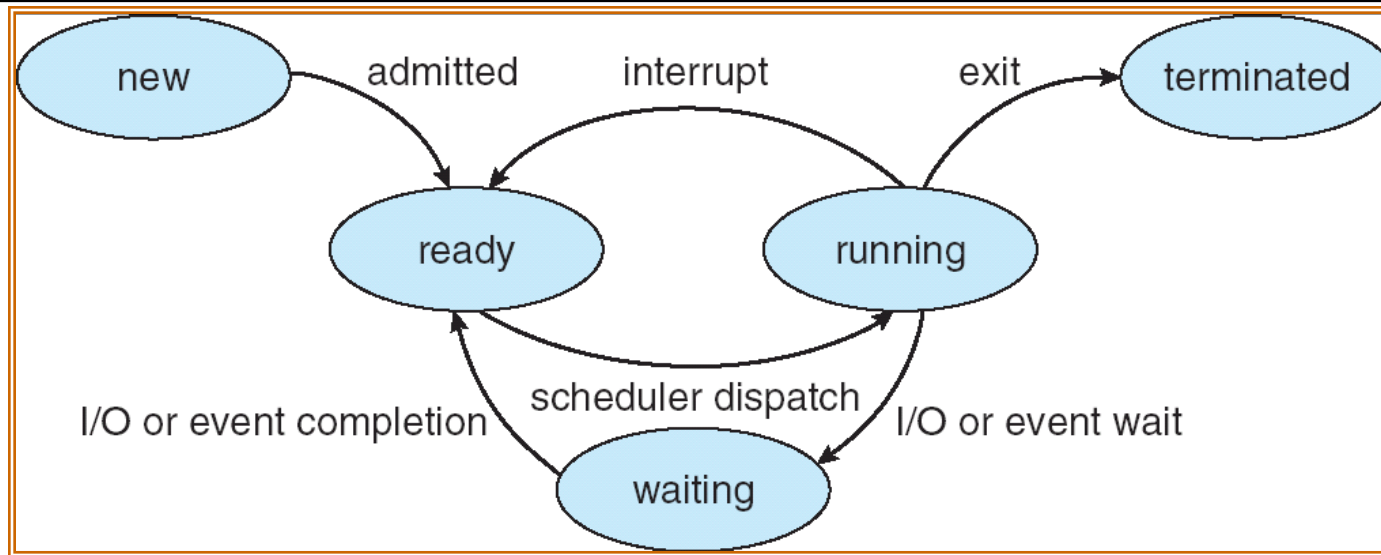
Process
Control
Block (PCB)

CPU Switch From Process to Process



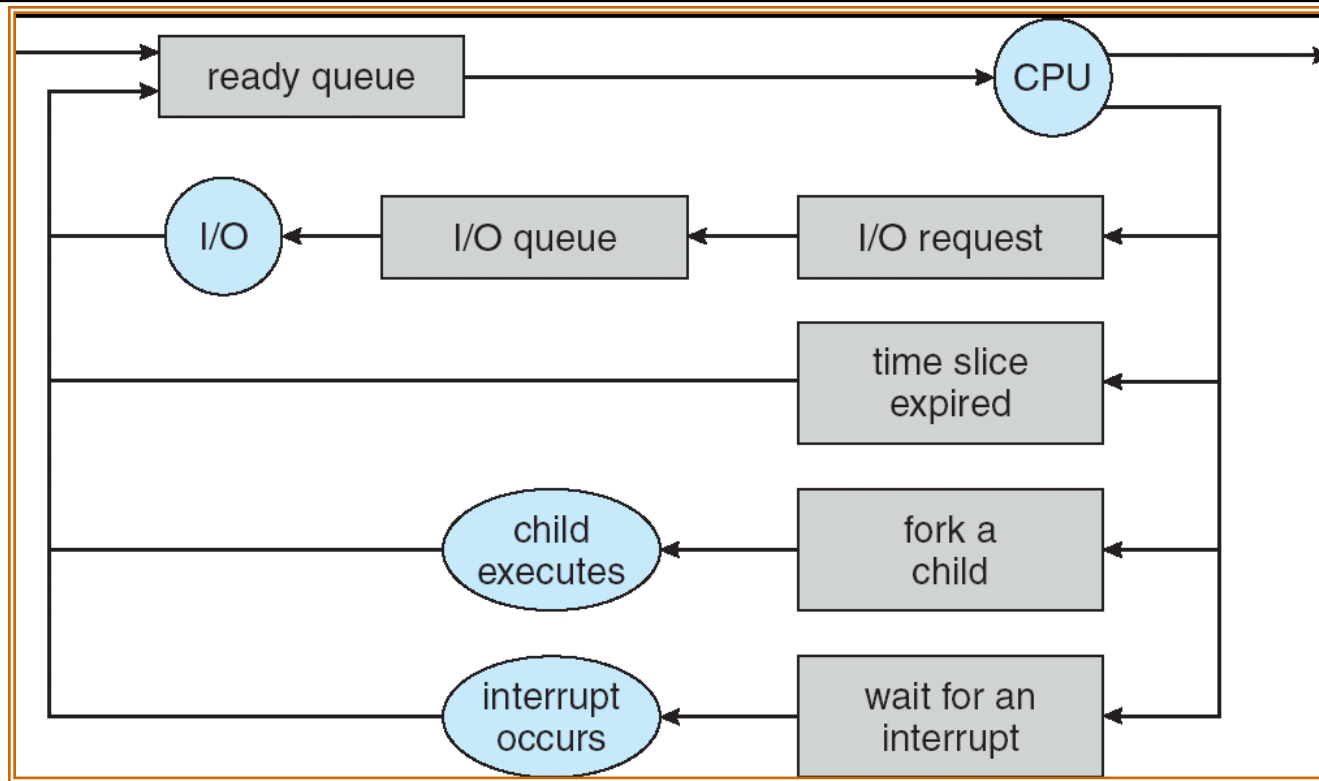
- This is also called a "context switch"
- Code executed in kernel above is overhead as the system does no useful work while switching
 - Overhead (\sim microsecs) sets minimum practical switching time
 - The more complex OS and PCB \rightarrow the longer context switch time
 - Some hardware provides multiple sets of registers \rightarrow multiple contexts loaded at once \rightarrow no saving/reloading overhead

Diagram of Process State



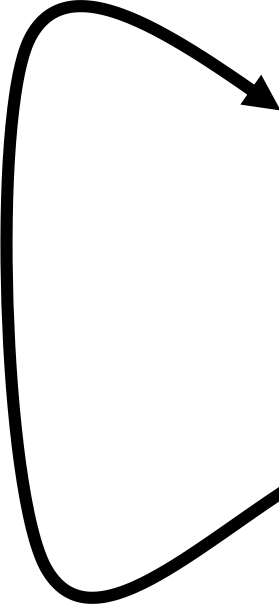
- As a process executes, it changes its state
 - **new**: The process is being created
 - **ready**: The process is waiting to run by CPU
 - **running**: Instructions are being executed by CPU
 - **waiting**: Process waiting for some event to occur
 - **terminated**: The process has finished execution

Process Scheduling (Queue representation)




- PCBs move from queue to queue as they change state
 - Decisions about which order to remove from queues are **Scheduling** decisions
 - Many algorithms possible for queues of CPU and I/O devices (discussed later)

CPU Scheduler

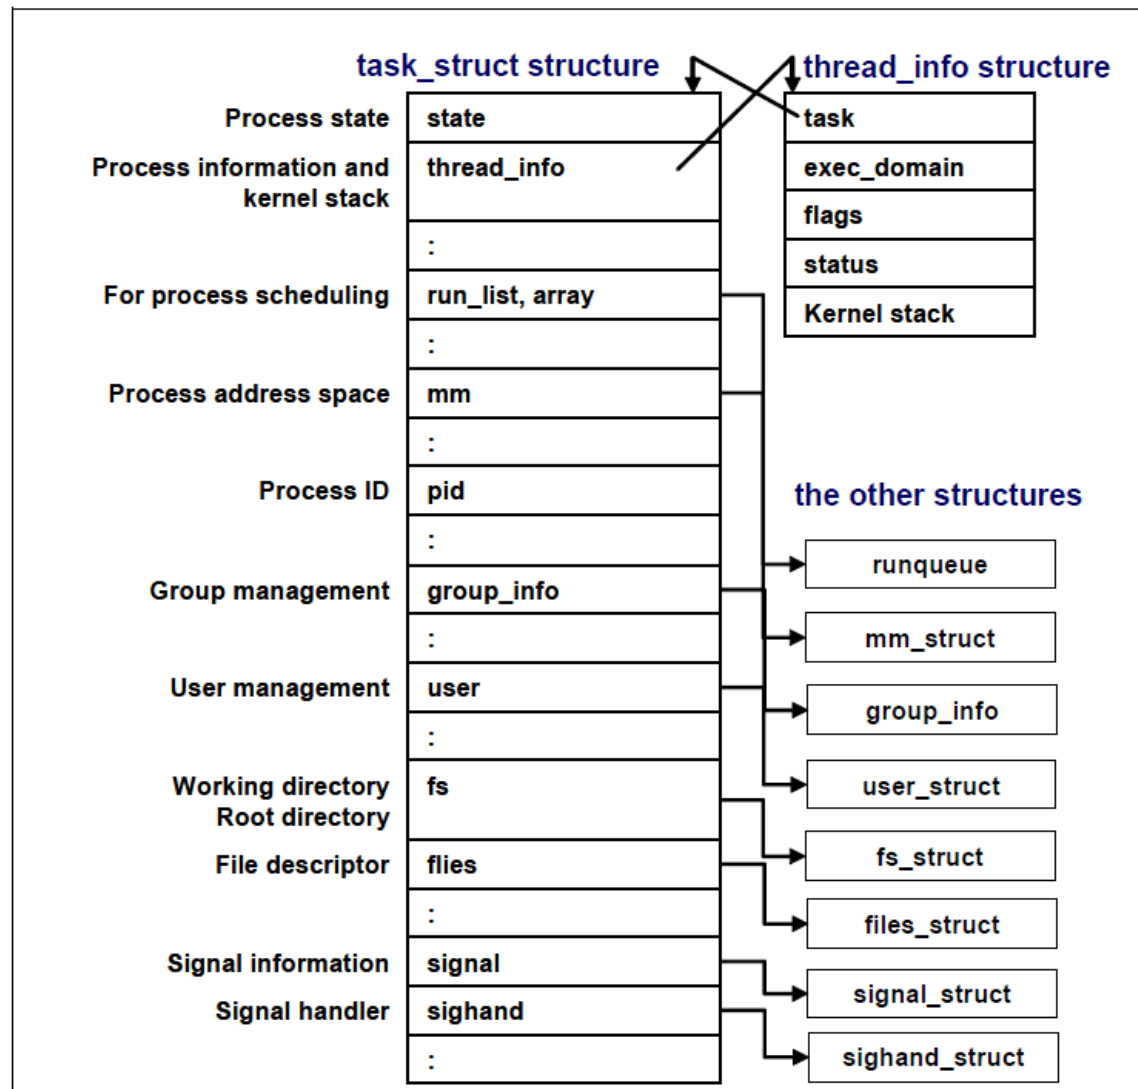


```
if ( readyProcesses(PCBs) ) {  
    nextPCB = selectProcess(PCBs);  
    run( nextPCB );  
} else {  
    run_idle_process();  
}
```

Ready Queue and I/O Device Queues



Linux Task Struct



Linux Task Struct (also called Process Descriptor):

<https://github.com/torvalds/linux/blob/master/include/linux/sched.h>

https://github.com/hungys/tech-note/blob/master/linux_kernel/process.md

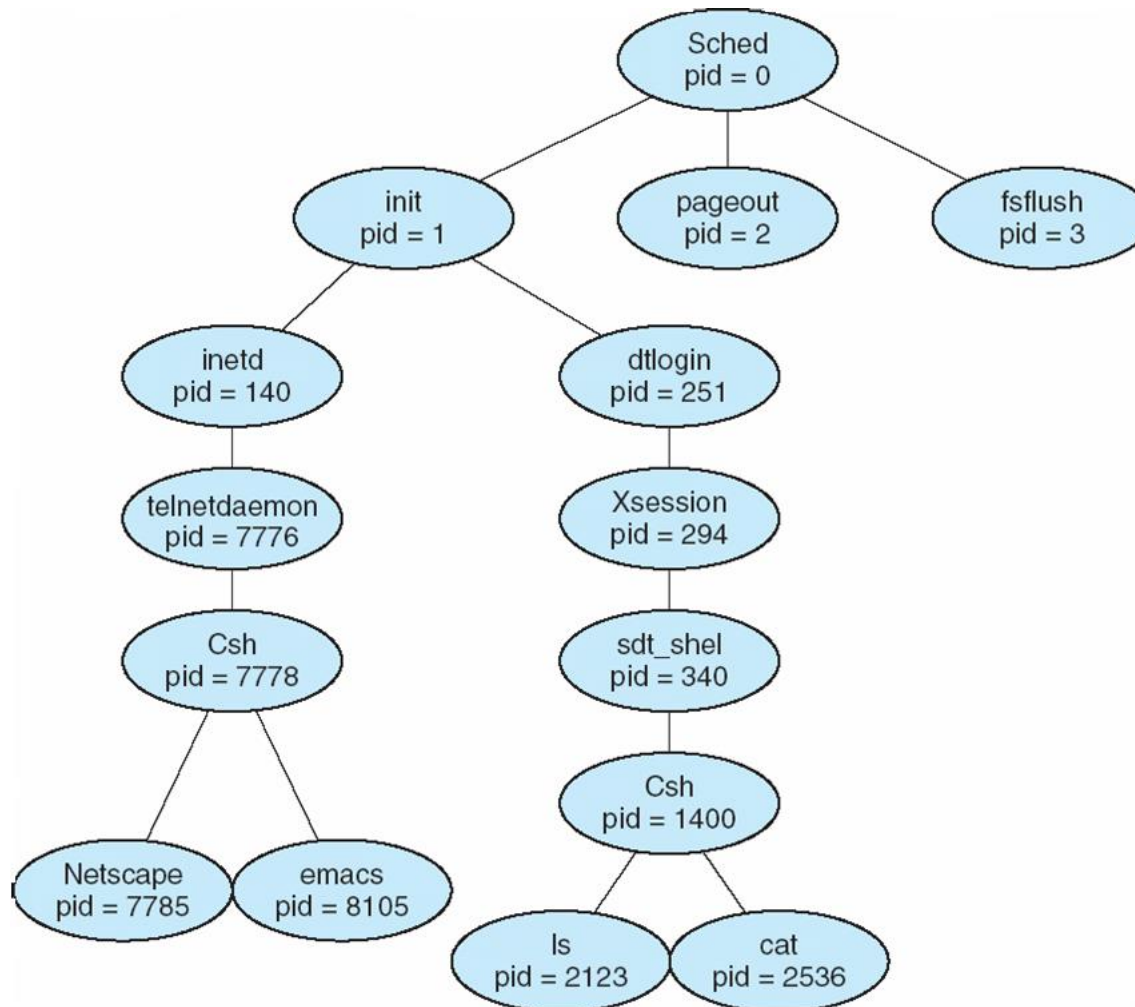
What does it take to create a process?

- Must construct new PCB
 - Inexpensive
- Must set up new page table for address space
 - More expensive (discussed later)
- Copy data from parent process? (Unix `fork()`)
 - Semantics of Unix `fork()` are that the child process gets a complete copy of the parent memory and I/O state
 - Originally very expensive
 - Much less expensive with "copy-on-write"
- Copy I/O state (file handles, etc)
 - Medium expense

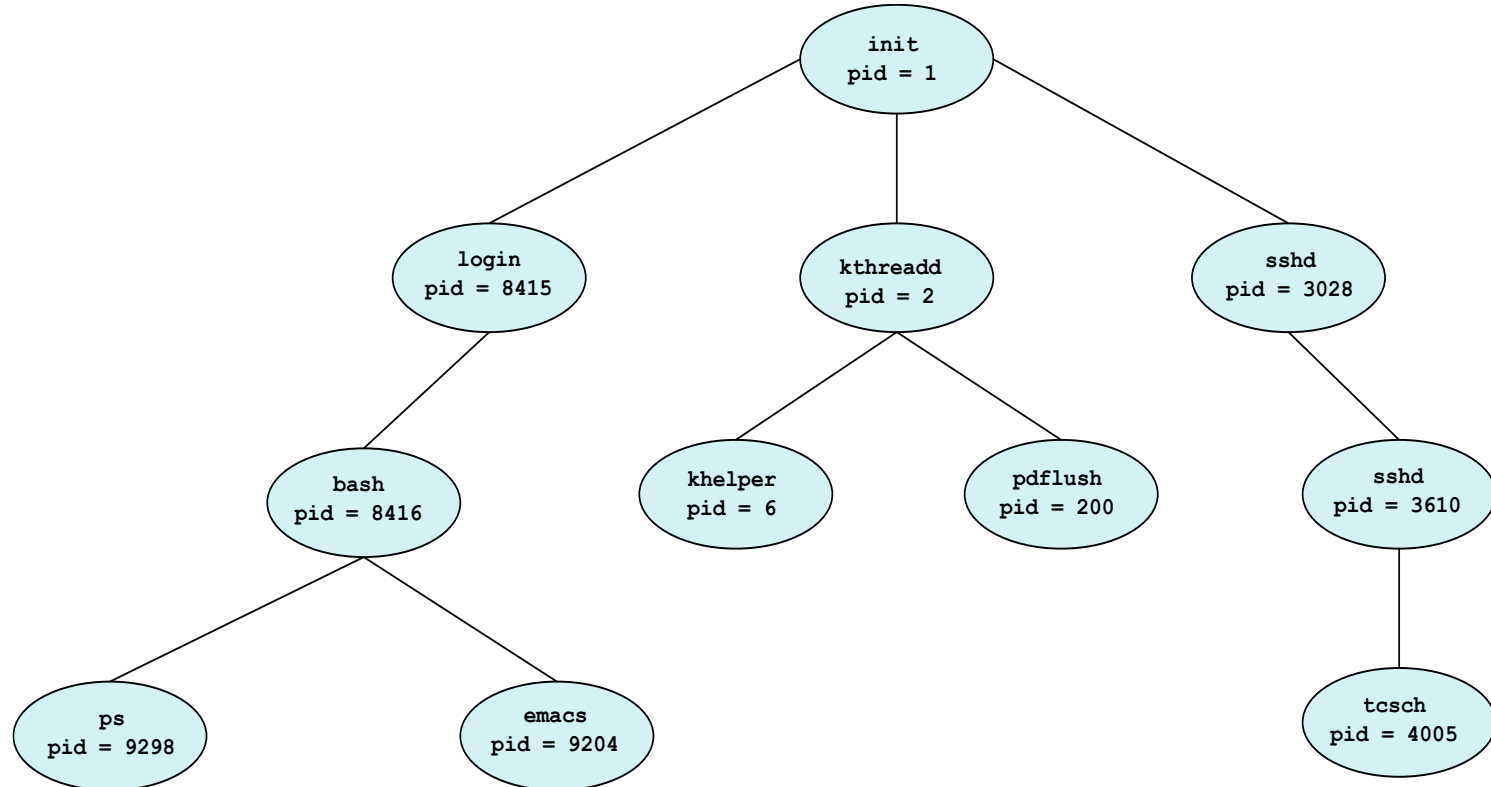
How to create a process?

- Double click on an icon or type program name on command-line (shell)
- After boot, OS starts the first process
 - E.g. sched for Solaris, ntoskrnl.exe for Windows
- The first process creates other processes:
 - the creator is called the parent process
 - the created is called the child process
 - parent/child relationships is expressed by a process tree
- Eg, in UNIX the second process is called *init*
 - it creates all the gettys (login processes) and daemons
 - it should never die
 - it controls the system configuration (#processes, priorities...)
- Explorer.exe in Windows for graphical interface

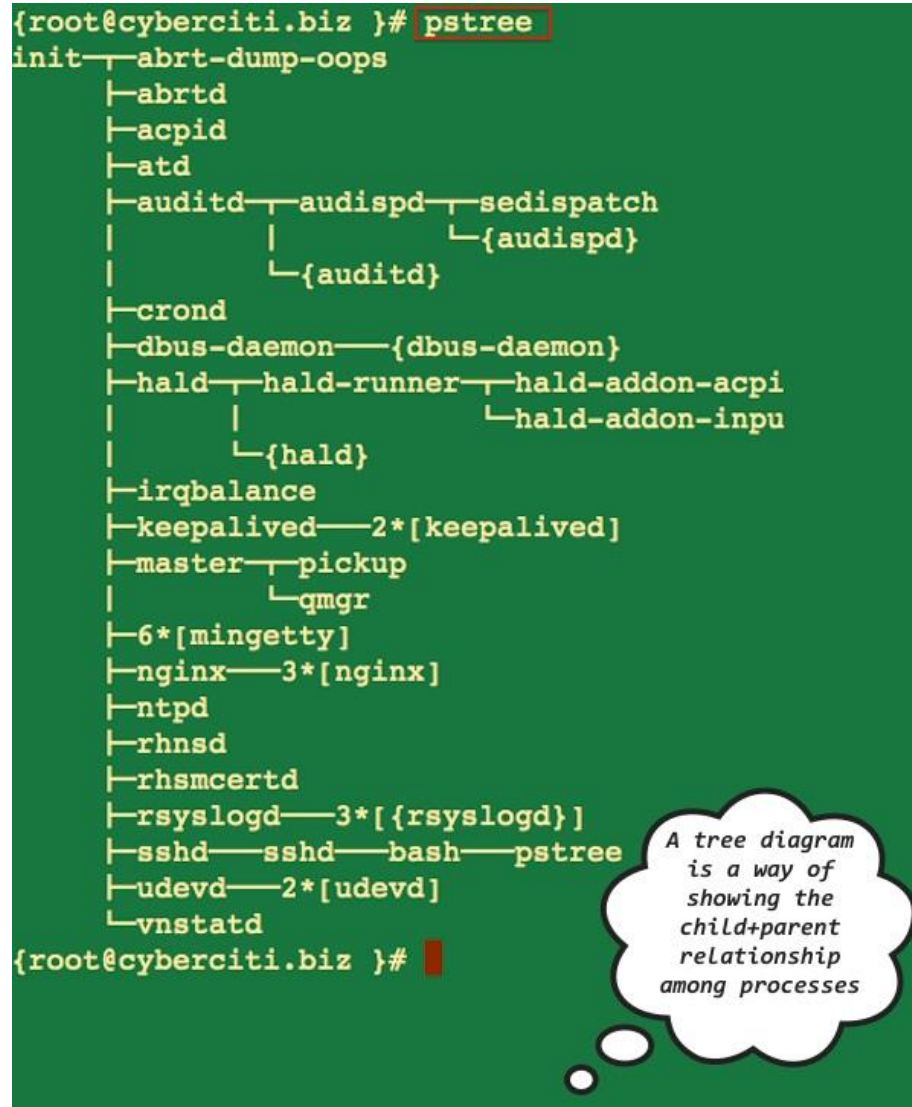
An Example Process Tree in Solaris



An Example Process Tree in Linux



An Example Process Tree in Solaris



Processes Under UNIX

- `Fork()` system call is only way to create a new process
- `int fork()` does many things at once:
 - creates a new address space (for the child)
 - copies the parent's address space into the child's
 - starts a new thread of control in the child's address space
 - parent and child are equivalent -- almost
 - » in parent, `fork()` returns a non-zero integer
 - » in child, `fork()` returns a zero.
 - » difference allows parent and child to distinguish
- `int fork()` returns TWICE!

ProcEx1.c: Program using fork()

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

ProcEx2.c

```
#include <stdio.h>
#include <sys/types.h> /* Primitive System Data Types */

int main(int argc, char **argv)
{
    char *myName = argv[0];
    int cpid = fork();
    if (cpid == 0) { //Child
        sleep(5); //sleeps for 5 secs
        printf("My pid: %d I am a child of %s My parent
pid: %d\n", getpid(), myName, getppid());
        exit(0);
    } else { //Parent
        printf("My pid: %d My child's pid: %d\n",
getpid(), cpid);
        exit(0);
    }
}
```

What does this program print?

ProcEx3.c

```
#include <stdio.h>
#include <sys/types.h> /* Primitive System Data Types */
#include <unistd.h>
#include <sys/wait.h>
int main(void)
{
    int i, pid, status;
    pid = getpid();
    fprintf(stdout, "parent pid = %d\n", pid);

    pid = fork();
    if (pid == 0) /* child process is always 0 */
    {
        for (i= 0; i < 10; ++i)
        {
            fprintf(stdout, "In child: Iteration: %d\n", i);
            sleep(0.1);
        }
        fprintf(stdout, "In child: child exiting\n");
    }
    else /* parent process is non-zero (child's pid) */
    {
        //sleep(2); //to force child to run first
        fprintf(stdout, "In Parent: child pid = %d\n", pid);
        fflush(stdout);
        fprintf(stdout, "In Parent: waiting for child\n");
        //wait(NULL); //wait for any child to change state
        //wait(&status); //status is stored here
        //waitpid(-1,&status,0); //wait for any child to change state
        waitpid(pid,&status,0); //wait for pid child to change state
        fprintf(stdout, "In Parent: Child exit status:%d\n",WEXITSTATUS(status));
        if(WIFEXITED(status))
            fprintf(stdout, "In Parent: Child exited normally\n");
        else if(WIFSIGNALED(status))
            fprintf(stdout, "In Parent: Child was killed by a signal!!\n");
        else
            fprintf(stdout, "In Parent: Child exited for other reasons\n");
        fprintf(stdout, "In Parent: child terminated\n");
        fprintf(stdout, "In Parent: parent exiting\n");
    }

    return 0;
}
```


Output of one of Runs: 1

parent pid = 26892
In Parent: child pid = 26893
In Parent: waiting for child
In child: Iteration: 0
In child: Iteration: 1
In child: Iteration: 2
In child: Iteration: 3
In child: Iteration: 4
In child: Iteration: 5
In child: Iteration: 6
In child: Iteration: 7
In child: Iteration: 8
In child: Iteration: 9
In child: child exiting
In Parent: Child exit status:0
In Parent: Child exited normally
In Parent: child terminated
In Parent: parent exiting

Output of one of Runs: 2

```
parent pid = 26898
In Parent: child pid = 26899
In child: Iteration: 0
In Parent: waiting for child
In child: Iteration: 1
In child: Iteration: 2
In child: Iteration: 3
In child: Iteration: 4
In child: Iteration: 5
In child: Iteration: 6
In child: Iteration: 7
In child: Iteration: 8
In child: Iteration: 9
In child: child exiting
In Parent: Child exit status:0
In Parent: Child exited normally
In Parent: child terminated
In Parent: parent exiting
```

ProcEx4.c

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>

int main(void)
{
    fprintf(stdout, "Parent PID: %d\n", getpid());
    fflush(stdout);

    while(1) {
        fork();
        fprintf(stdout, "My PID: %d and My Parent
                                PID: %d\n", getpid(), getppid());
    }
    return 0; }
```

ProcEx5.c

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>

int main(void)
{
    fprintf(stdout, "Parent PID: %d\n", getpid());
    fflush(stdout);

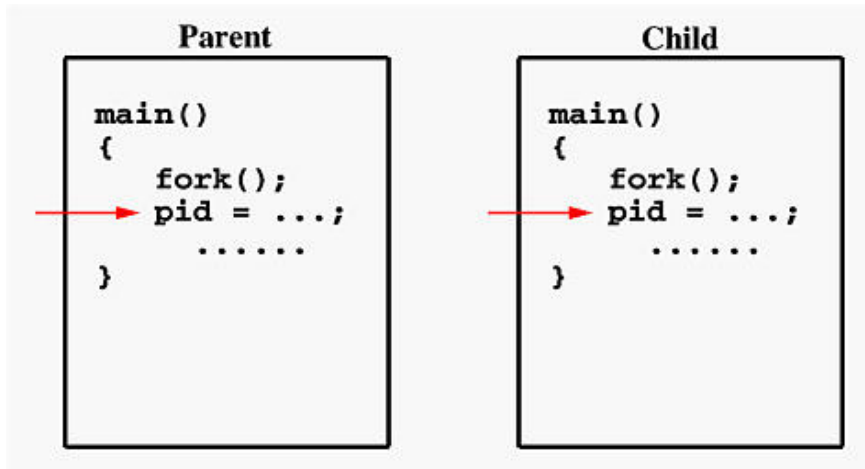
    while(fork()) wait(NULL);
    fprintf(stdout, "My PID: %d and My Parent
                                   PID: %d\n", getpid(), getppid());
    return 0;
}
```

More Examples on fork()

- <http://www.csl.mtu.edu/cs4411.ck/www/NOTES/process/fork/create.html>
- <http://www.amparo.net/ce155/fork-ex.html>
- http://home.adelphi.edu/sbloch/class/archive/271/fall2005/examples/c/fork_examples/
- <http://man7.org/linux/man-pages/man2/vfork.2.html>
 - Creates a child process and block parent till child finishes
 - Till finishes the child shares all memory with its parent, including the stack!

Fork is half the story

- Fork() gets us a new address space for child,
 - but parent and child share EVERYTHING
 - » memory, operating system state



- int **exec**(char ***prgName**) completes the picture
 - throws away the contents of the calling address space
 - replaces it with the program named by **prgName**
 - starts executing at header.startPC
 - **exec** does not return on successful load of new prog!
 - Returns -1 if it fails to load new prog
- Pros: Clean, simple; Con: duplicate operations


exec family

- `execi("/bin/vi","vi","/home/user/ex1.txt",NULL); //PATH`
- `execip("vi","vi","/home/user/ex1.txt",NULL); //File`
- `execle(path, arg1,..., envp[])`
- `execv(path, argv[]), execvp(progName, argv[])`
- `execvpe(progName, argv[], envp[])`
- **i**: list of Args, **v**: vector of Args
- **p**: Searches folders listed in PATH environment variable
 - `$echo $PATH` or `$env |grep PATH` → `/usr/bin:/usr/local/bin`
- **e**: environment of progName to be specified
- The first argument should point to the PrgName associated with the file being executed.
- The list of arguments *must* be terminated by a null pointer
- `execip()`, `execvp()`, and `execvpe()` **duplicate actions of the shell** in searching for an executable file if the specified PrgName does not contain / character
- Programs executed from the shell inherit all of the environment variables from the shell.

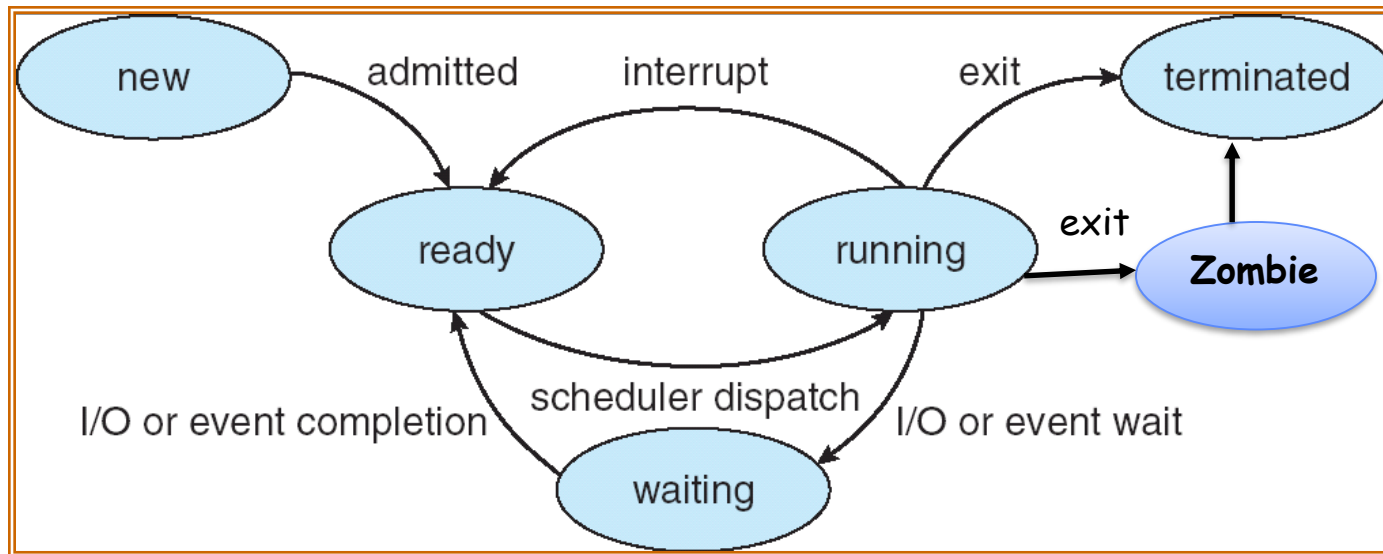
Process Termination

- Process executes last statement and OS decides(exit)
 - Output/return data from child to parent (via wait)
 - Process' resources are deallocated by operating system
- Child exits before parent calls up waitpid()
 - Child is in **Zombie state** and has entry still maintained in PCB linked list
 - It is removed later when parent calls waitpid()
- Parent may terminate execution of child process (abort)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting first, child becomes orphan (**it's not a state**)
 - » Some OSes don't allow child to continue if parent terminates
 - All children terminated - *cascading termination*
 - Linux: **init daemon** becomes parent of these **orphan children**; **upstart** - init replacement daemon in Ubuntu; **systemd** - init replacement daemon designed to start processes in parallel → quick booting of OS

Process Life Cycle



Process State Diagram



Multiprocess Architecture - Chrome Browser

- Many web browsers ran as single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - **Browser** process manages UI, disk and network I/O
 - **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened on a tab
 - » Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in



Multitasking in Mobile Systems

- Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- Due to screen real estate & user interface limits, iOS provides for a
 - Single **foreground** process- controlled via user interface
 - Multiple **background** processes- in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
 - Background process uses a **service** to perform tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use

Reading Assignment

- Chapter 3 from OSC by Galvin et al
- Prof. John Kubi Video Lectures: [L1](#) and [L2](#)
- Chapter 2 from OSD&I by Tanenbaum et al
- [Chapter 3. Processes](#) and [Chapter 19. Process Communication](#) from Understanding the Linux Kernel by Daniel P. Bovet and Marco Cesati (available on Intranet)
- <http://man7.org/linux/man-pages/man2/kill.2.html>
- https://www.gnu.org/software/libc/manual/html_node/index.html
- <https://www.chromium.org/developers/design-documents/process-models>
- <http://www.tldp.org/LDP/Linux-Filesystem-Hierarchy/html/proc.html>
 - time cmd
 - /usr/bin/procinfo